Methods: SkillTree Export, Dashboard Ownership, and Publication Surfaces

SkillTree export as a dependency-safe project, subject, skill, and quiz map I

`src/skilltree` converts the validated, provenance-rendered graph into SkillTree's Project -> Subject -> Skill model, emitting prerequisite **learning-path dependencies** and a **quiz** per skill from its assessment questions. A REST client and seeder create the project on a configured SkillTree instance, when available, in dependency-safe topological order: subjects, then quizzes, then skills, then dependencies. Quiz definitions must exist before quiz-gated skills are created, and dependencies are assigned only after both endpoint skills exist. The REST contract is verified against a local `pytest-httpserver` mirror of SkillTree's documented endpoints, so the integration is tested without requiring a running Java service.

Learner dashboard as a local lesson surface and read-only SkillTree handoff I

The local learner dashboard is intentionally a shell rather than a competing learning runtime. It reads generated artifacts (`skilltree_project.json`, claims, stats, figures, and the rendered manuscript), probes SkillTree's public status endpoint, and mounts the official `@skilltree/skills-client-js` learner display when a local trusted-client bridge is ready. The bridge keeps SkillTree credentials server-side: browser code sees only the configured service URL, project id, and a dashboard-local token endpoint.

The curriculum explorer in the dashboard is therefore read-only. It renders the exported subjects, skills, quiz identifiers, and prerequisites so authors can inspect the exact artifact that will be seeded. The **Learn** action opens a local modular lesson record for the selected skill from `course_material_index.json`: concept, prerequisite context, downstream role, assessment summary, provenance note, related local skills, and pinned external pointers. **Open Subject** exits to the official SkillTree subject learning path (`/subjects/{subjectId}`), where SkillTree lists subject progress, rank, and quiz affordances. **Quiz** deep-links to `/subjects/{subjectId}/skills/{skillId}/quizzes/{quizId}`. This distinction matters operationally. Lesson summaries, search, artifact previews, and status summaries are project-owned conveniences; progress, ranking, scoring, and completion remain SkillTree-owned

Learner dashboard as a local lesson surface and read-only SkillTree handoff I

learner state.

Publication surfaces for cover, page-two atlas, figures, and artifact audits I

The first and second pages are also generated, not hand-composed. `scripts/run_kernels.py` writes `active_skillference_cover.png` from the live curriculum graph, claim registry and roadmap metrics, and writes `active_skillference_skill_atlas.png` from every authored skill name in the curriculum graph. The cover audit checks that `book.cover.image` resolves to the cover PNG; the front-matter visual audit checks that `front_matter.page_two_visual.image` resolves to the atlas PNG. Both audits require a nonblank image, dashboard exposure as non-figure artifacts, and exclusion from `figure_registry.json`. This keeps the public entry surfaces aligned with the same artifact layer that drives the results figures while preserving the stricter cross-reference contract for numbered

Publication surfaces for cover, page-two atlas, figures, and artifact audits I

manuscript figures.