Scope: Curriculum Platform, Runtime Ownership, and Non-Claims

Curriculum platform scope: not an inference engine or hosted LMS I

Active Skillference is curriculum-authoring and delivery infrastructure, not a new inference engine. Its kernels are pedagogical, correctness-first reference implementations chosen for legibility and testability, not for the performance or breadth of production toolboxes such as `pymdp` for discrete state spaces, concise discrete-time tutorial/code treatments aligned with that ecosystem (Heins et al. 2022; Oostrum et al. 2025), or SPM/DEM for continuous dynamic systems (Friston et al. 2007, 2008). The contribution described here is *how the curriculum is built and verified*: a validated prerequisite DAG with content-provenance binding, exported into an existing micro-learning runtime. The repository is not an intelligent tutor, a learner-state estimator, or a general-purpose adaptive system.

Curriculum platform scope: not an inference engine or hosted LMS I

The curriculum is deliberately broad (630 skills across all eight strata) and designed so that growth is a data operation: adding skills or worked examples extends the catalog without changing the engine. Its new practice layer is still pedagogical rather than clinical or operational advice: domain-fit, ethics, misuse, and evaluation skills train learners to ask better modeling questions, not to deploy autonomous interventions without local review. Out of scope for this release are hosting and authentication of a production SkillTree deployment, learner analytics, and a textbook-length treatment of every individual skill. The local dashboard narrows that boundary: it is enough to show that the export can be inspected locally and, when a SkillTree runtime is available, opened through subject and quiz routes; it is not a hosted LMS or analytics product.