Code and Data Availability

Code and Data Availability I

The artifact described here is a single repository containing the kernels, the curriculum catalog, the SkillTree exporter, the learner dashboard, the manuscript sources, and every validator named in sec. 14. Authored code and content are MIT licensed with Daniel Ari Friedman as the copyright holder, recorded in `LICENSE` and mirrored in the citation, CodeMeta, and Zenodo metadata files. Vendored runtime dependencies and their separate notices are listed in `THIRD_PARTY_NOTICES.md`. External Active Inference Institute course and knowledge-base material is referenced by repository, path, and commit rather than copied, so upstream licensing stays with upstream.

Code and Data Availability I

The code is public at github.com/ActiveInferenceInstitute/Active_Skillference, with the v1.0.0 manuscript archived on Zenodo at doi.org/10.5281/zenodo.21865644. The identifier appears on the cover page, in the citation file, in the CodeMeta and Zenodo records, and in the archived copy of this document, all naming the same deposit. That agreement is enforced rather than assumed: a test cross-checks each surface against the recorded deposit and fails on any divergence, which is what keeps a plausible but wrong identifier out of a citation. Version-specific identifiers stay blank because no versioned deposit beyond this one exists, and the same test holds them blank.

Code and Data Availability I

Generated outputs travel with the repository even though they are rebuildable, so a reader can open the manuscript, the figures, the SkillTree export, and every audit record without installing anything or running the pipeline. Committing them does not soften the regenerability claim: running the kernels rebuilds the claim ledger, manuscript variables, scholarship audits, figures, lesson records, and the export from source, and the freshness validator fails if a committed artifact has drifted from what the current source produces. Only build detritus is withheld, meaning renderer logs and intermediates that change on every run and carry nothing a reader would read. The captured evidence described in sec. 11 is different in kind rather than in handling: a recording of a running system is not a function of project data and cannot be rebuilt from a checkout at all, which is why it is stored beside a provenance record

Code and Data Availability I

naming the dashboard address, viewport, live counts, full-stack mode, capture timestamp, and a digest per image. Every other image in this paper regenerates; those five are archived precisely because they do not.